

DATASCI207-005/007

Applied Machine Learning

Vilena Livinsky, PhD(c)

School of Information, UC Berkeley

Week 2: 01/15/2025 & 01/16/2025

Today's Agenda

- Logistics note: Final Project
- Linear Regression & Gradient Descent
- Walkthroughs:
 - Gradient Descent & Linear Regression
 - TensorFlow Intro



Final Project: Step 1



Form a group

3-4 people, max 4

NO silos or groups of 2

Groups can only be composed of you and your colleagues in your section



Inform me & the class of your formed group in Slack

Include names of group members

Due date: 01/24/2025 EOD



General Plan:

Step 1: form a group

Step 2: submit your group's question to investigate + dataset

Step 3: a baseline presentation

Step 4: final presentation

Walkthrough

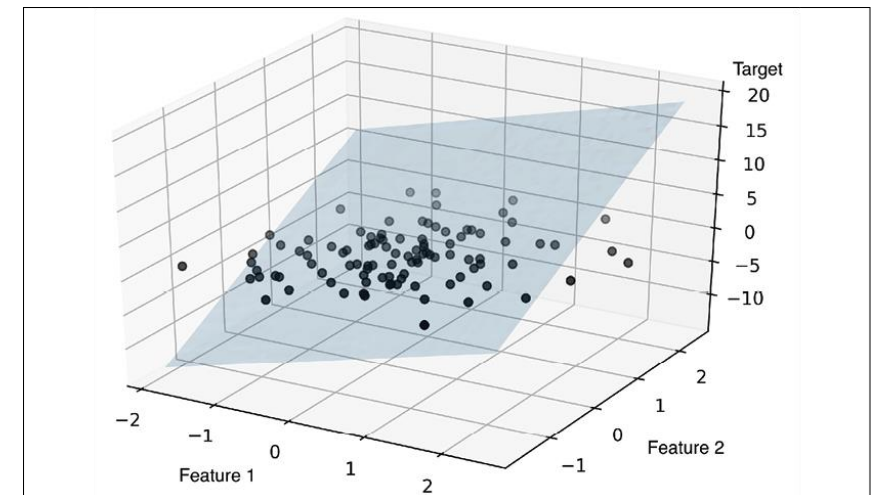
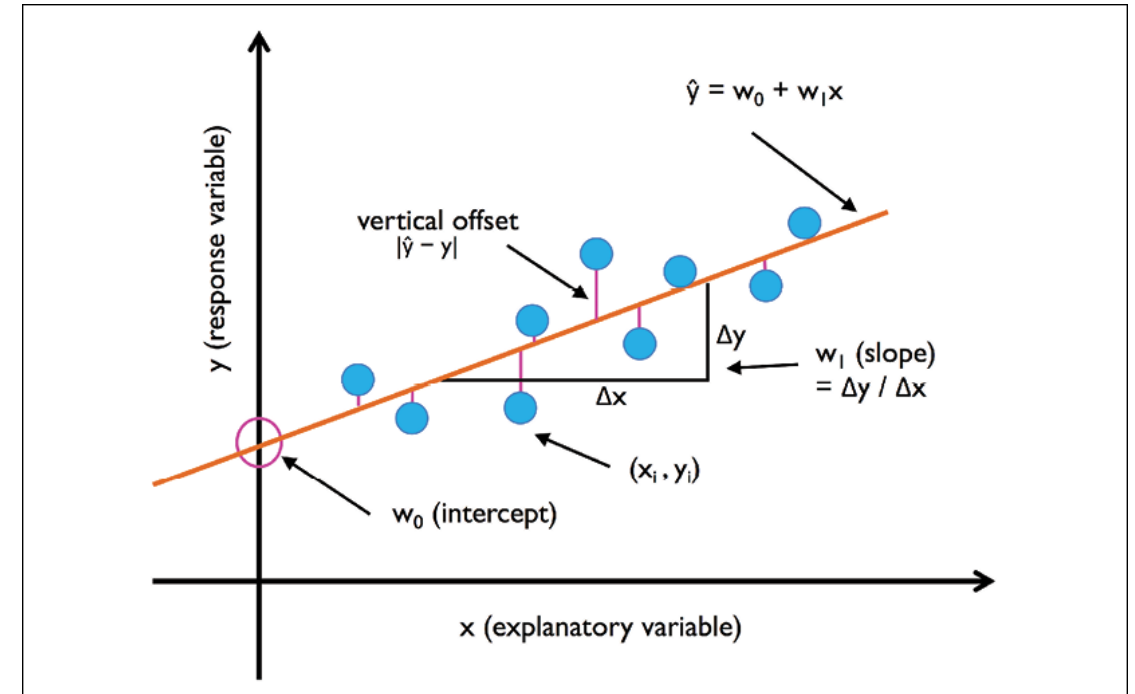
01b_Framing.ipynb

Linear Regression & Gradient Descent (Review)

- Supervised learning
- predict target variable on a continuous scale
- simple (univariate) linear regression vs. multivariate case

$$y = w_0x_0 + w_1x_1 + \dots + w_mx_m = \sum_{i=0}^n w_ix_i = w^T x$$

Here, w_0 is the y axis intercept with $x_0 = 1$.



Linear Regression & Gradient Descent (Review)

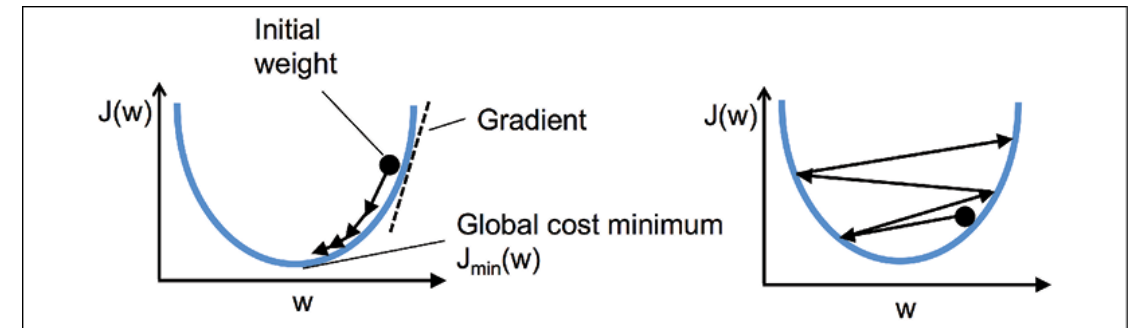
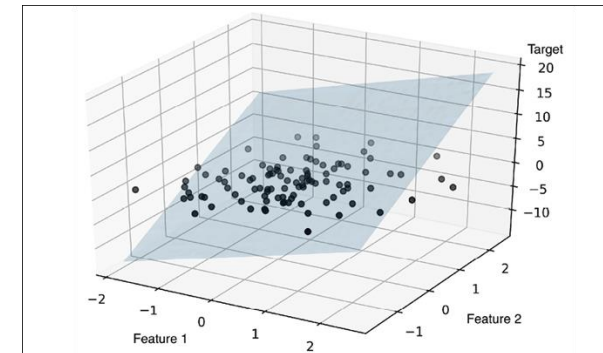
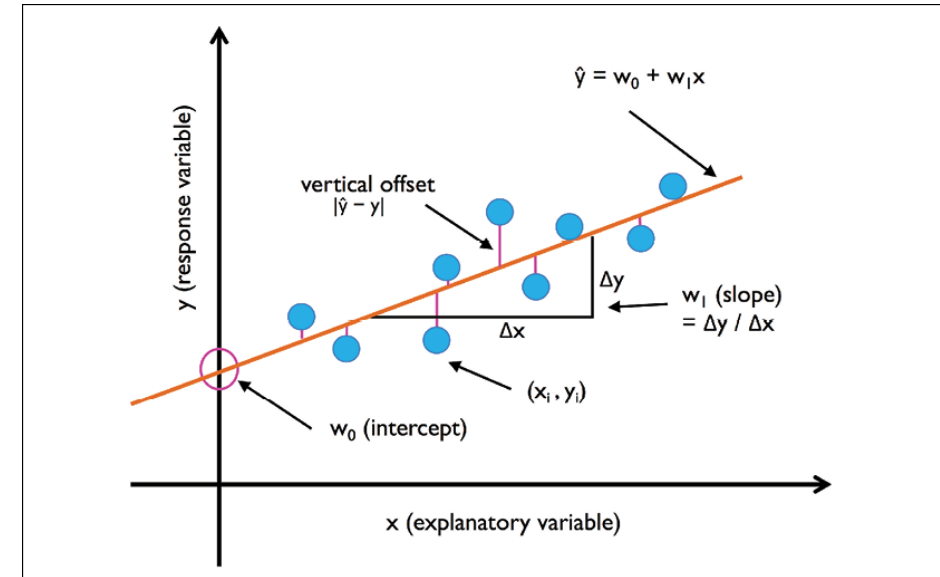
- Supervised learning
- predict target variable on a continuous scale
- simple (univariate) linear regression vs. multivariate case
- Gradient Descent:
 - Choose some initial value for \mathbf{w}
 - Should we increase or decrease w ?
 - Use the slope (gradient)
 - Keep updating \mathbf{w} by following the gradient, until we reach convergence
 - gradient

update (w) rule

$$w := w - \alpha \frac{\partial}{\partial w} J$$

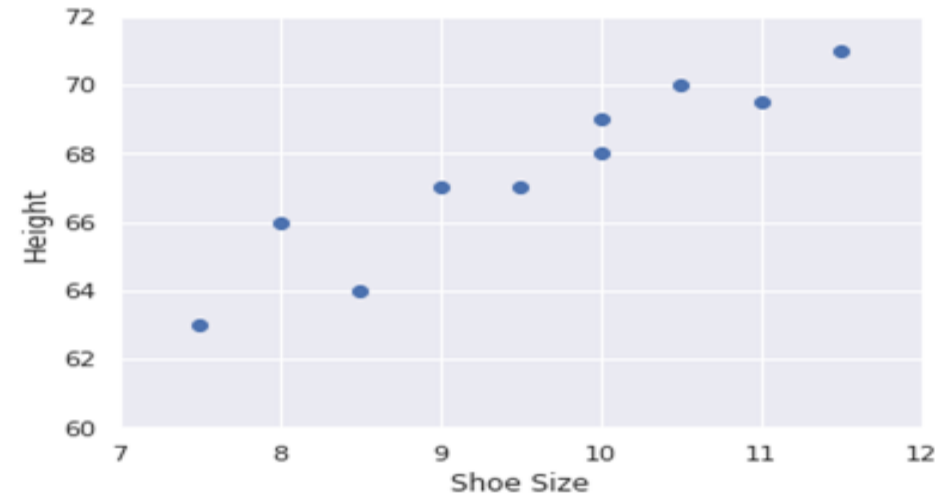
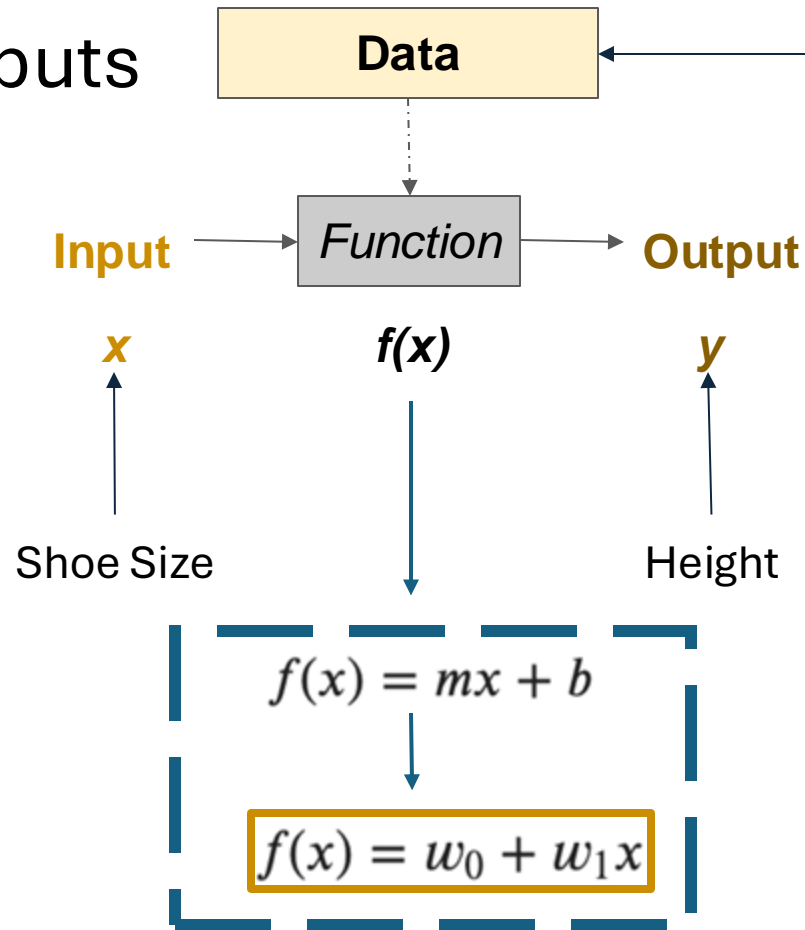
Ex: Linear Regression: MSE

$$\frac{\partial}{\partial w} J \rightarrow \frac{\partial}{\partial w_1} J = \frac{\partial}{\partial w_1} (\hat{y} - y)^2 \rightarrow 2(\hat{y} - y)x$$



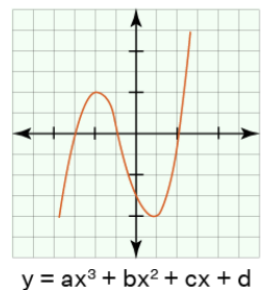
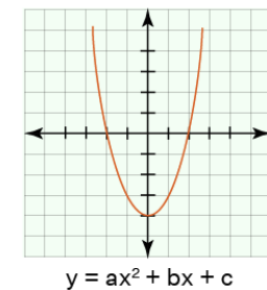
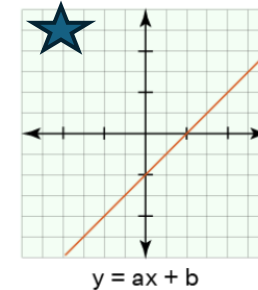
ML Framing: Linear Regression

- Inputs and Outputs
- Labeled data
- Split train/test
- Evaluation
- Baseline
- Use in practice



Graphs of Polynomial Functions

<https://www.cuemath.com/algebra/linear-polynomial/>



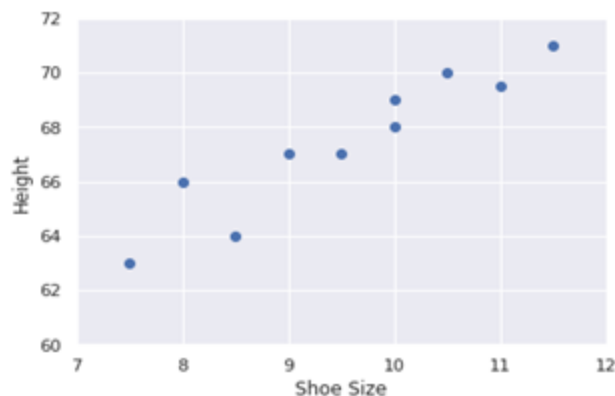
Review: The One-Variable Case

Model $f(x) = w_0 + w_1 x$

Parameters $W = [w_0, w_1]$

Loss $J(w_0, w_1) = \frac{1}{m} \sum_i (w_0 + w_1 x_i - y_i)^2$

Objective *Minimize* $J(w_0, w_1)$



$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Update Rule(s) for w

$$w := w - \alpha \frac{\partial}{\partial w} J$$

$$\begin{cases} w_0 := w_0 - \alpha \frac{1}{m} \sum_i (w_0 + w_1 x_i - y_i) \\ w_1 := w_1 - \alpha \frac{1}{m} \sum_i (w_0 + w_1 x_i - y_i) x_i \end{cases}$$

Dot Product & Matrix Transpose

- Sum of the products of the values in \mathbf{x} and \mathbf{w} using a (vector) dot product:

$$z = w_0x_0 + w_1x_1 + \dots + w_mx_m = \sum_{j=0}^m x_jw_j = \mathbf{w}^T\mathbf{x}$$

$$\begin{bmatrix} 1 & 2 & 3 \end{bmatrix} \times \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} = 1 \times 4 + 2 \times 5 + 3 \times 6 = 32$$

- Transpose (T) is only applicable to matrices
 - In ML:
 - $n \times 1$ or $1 \times m$ matrices are referred to as vectors

- Matrix multiplication ~ vector dot product (YAY NumPy!)

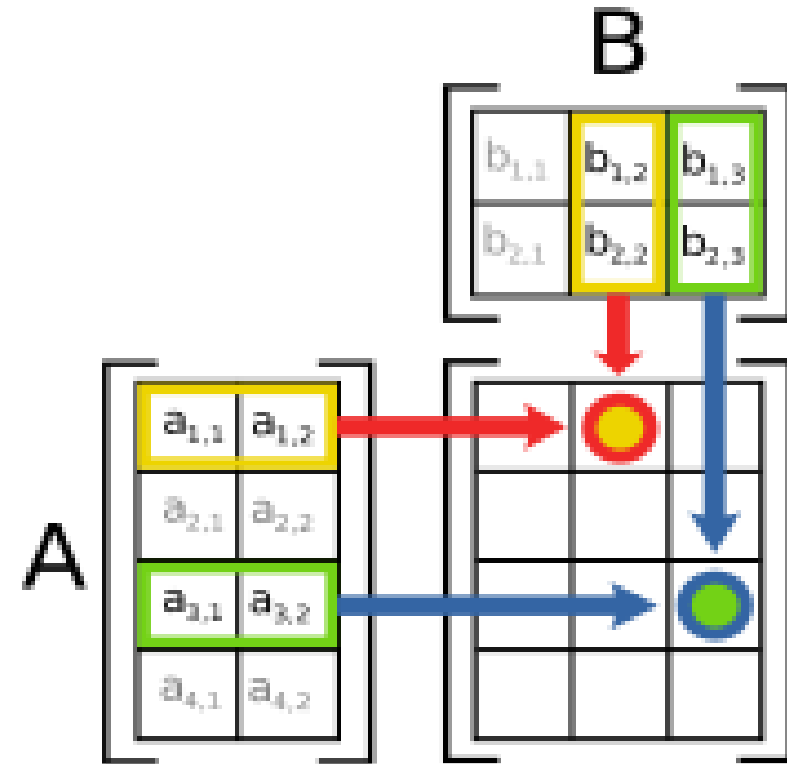
- Each row in the matrix treated as a single row vector
- Efficient computation

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 1 \times 7 + 2 \times 8 + 3 \times 9 \\ 4 \times 7 + 5 \times 8 + 6 \times 9 \end{bmatrix} = \begin{bmatrix} 50 \\ 122 \end{bmatrix}$$

Matrix Multiplication: Review

$$\begin{array}{c}
 \text{4} \times \text{2 matrix} \\
 \begin{bmatrix} a_{11} & a_{12} \\ \cdot & \cdot \\ a_{31} & a_{32} \\ \cdot & \cdot \end{bmatrix}
 \end{array}
 \begin{array}{c}
 \text{2} \times \text{3 matrix} \\
 \begin{bmatrix} \cdot & b_{12} & b_{13} \\ \cdot & b_{22} & b_{23} \end{bmatrix}
 \end{array}
 =
 \begin{array}{c}
 \text{4} \times \text{3 matrix} \\
 \begin{bmatrix} \cdot & c_{12} & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & c_{33} \\ \cdot & \cdot & \cdot \end{bmatrix}
 \end{array}$$

$$\begin{aligned}
 c_{12} &= a_{11}b_{12} + a_{12}b_{22} \\
 c_{33} &= a_{31}b_{13} + a_{32}b_{23}.
 \end{aligned}$$



Linear Regression & Gradient Descent (Review)

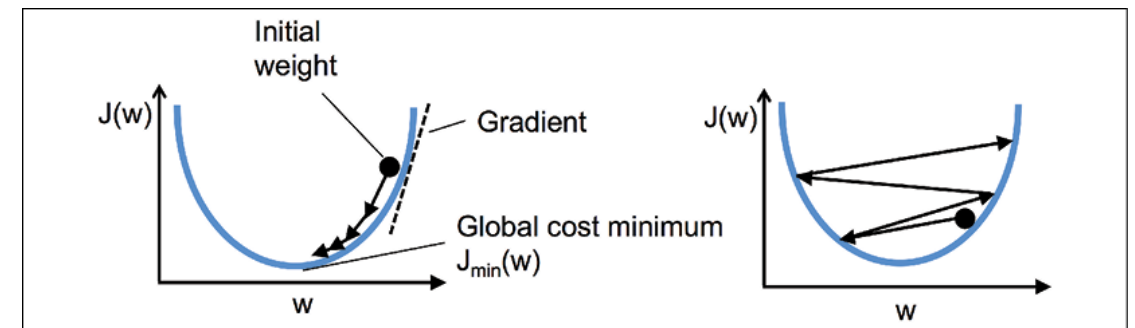
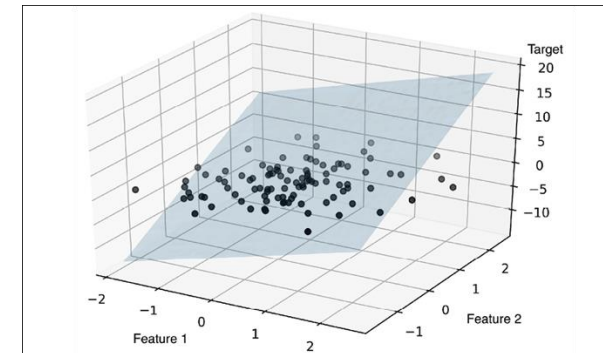
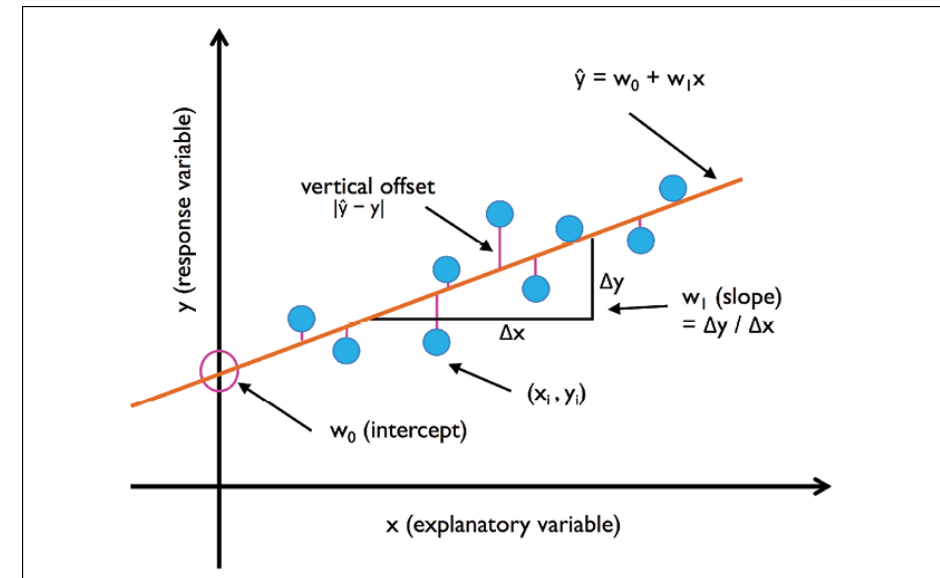
- Training:
 - Get data
 - Initialize weight/s
 - Initialize bias
- Given data
 - Get prediction
 - Get error
 - Update weight, use Gradient Descent
 - Repeat

```
# Run gradient descent
learning_rate = 0.1

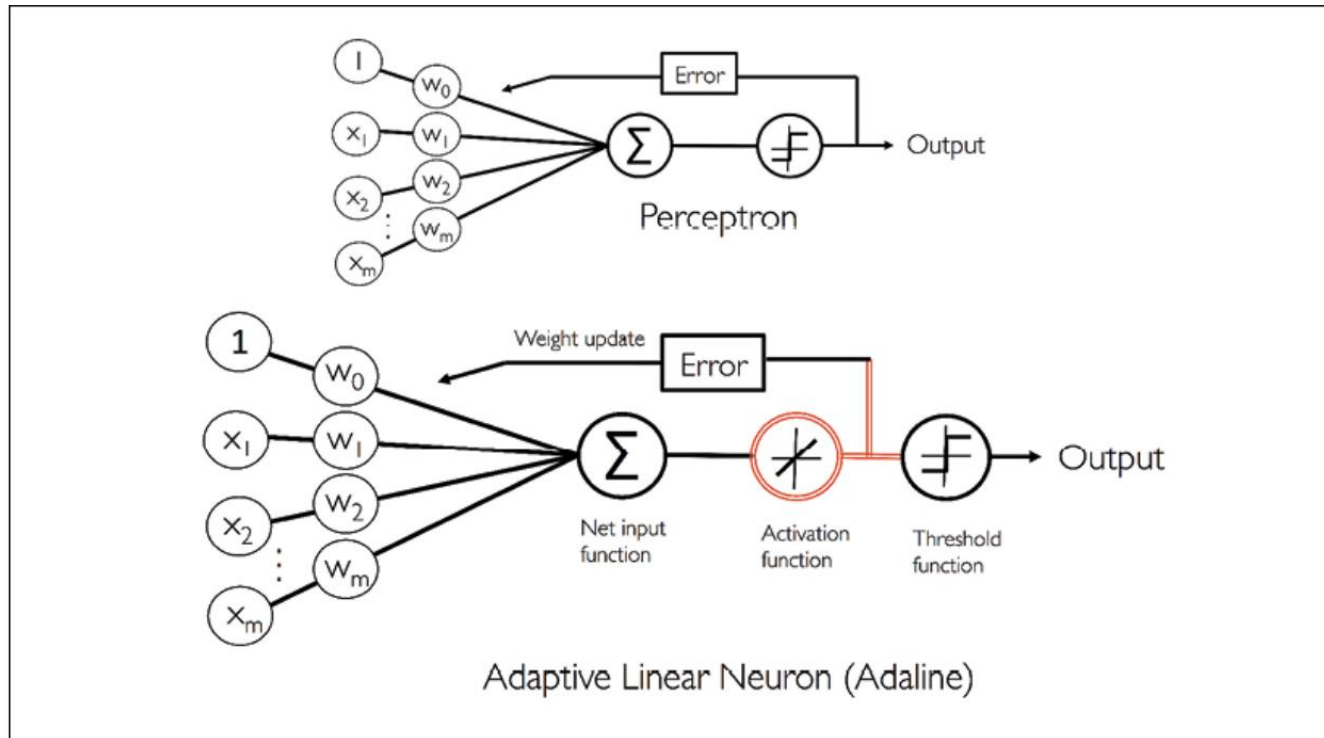
preds = np.dot(X, W)
loss = ((preds - Y)**2).mean()
gradient = 2 * np.dot((preds - Y), X) / m
W = W - learning_rate * gradient

print('predictions:', preds)
print('loss:', loss)
print('gradient:', gradient)
print('weights:', W)

predictions: [6 5]
loss: 16.0
gradient: [ 8. 20. 12.  4.]
weights: [ 0.2 -1. -0.2  0.6]
```



Single Layer Neural Net & Activation Function(s)



net input $z = \mathbf{w}^T \mathbf{x}$

↓

linear activation function, $\phi(z)$,

$$\phi(\mathbf{w}^T \mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

Appendix

Extra slides for review

Closed Form Solution for Solving OLS

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

```
# adding a column vector of "ones"
>>> Xb = np.hstack((np.ones((X.shape[0], 1)), X))
>>> w = np.zeros(X.shape[1])
>>> z = np.linalg.inv(np.dot(Xb.T, Xb))
>>> w = np.dot(z, np.dot(Xb.T, y))
>>> print('Slope: %.3f' % w[1])
Slope: 9.102
>>> print('Intercept: %.3f' % w[0])
Intercept: -34.671
```