

DATASCI207-005/007

Applied Machine Learning

Vilena Livinsky, PhD(c)

School of Information, UC Berkeley

Week 5: 02/05/2025 & 02/06/2025

Today's Agenda

- Multiclass Classification & Metrics
- Walkthroughs:
 - Metrics
 - Multiclass Classification
 - + TensorFlow



Practice



TensorFlow: General Modeling Steps

Create

Create a model

- Create your model architecture
- *Functional* or *Sequential API*

```
# Use Keras Sequential API to build a linear regression model.
model = keras.Sequential()
model.add(keras.layers.Dense(
    input_shape=[num_features], # each input has num_features features
    units=1,                    # there is a single output
    use_bias=True               # include a learned bias parameter
))
```

Compile

Compile a model

- Decide on **loss**
- Decide on model performance metrics
- Decide on how to improve the process; choose an **optimizer**

```
model.compile(loss='mean_squared_error', optimizer='adam',
              metrics=['mean_absolute_error'])
```

loss vs. metrics?

Fit

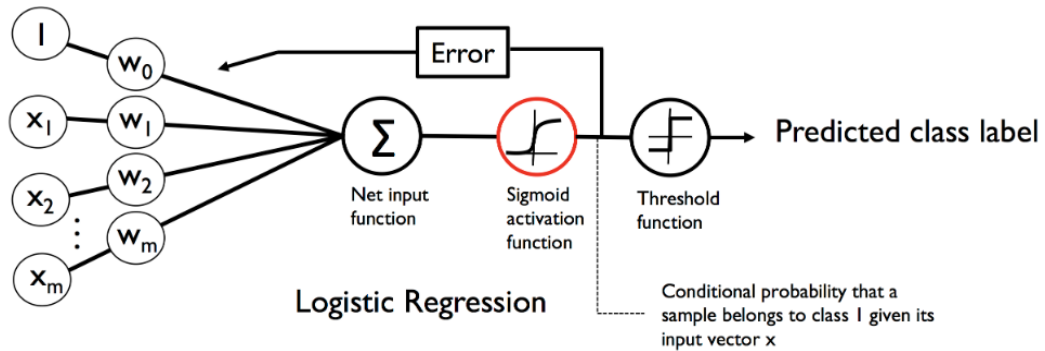
Fit a model

- Learn a function

```
# Build a model and train it. Hold out 10% of data for validation.
model = build_model(num_features=len(features))
model.fit(x=X, y=Y,
          validation_split=0.1, batch_size=16, epochs=5)

# Use the model to predict the training labels.
Y_pred = model.predict(x=X).flatten()
```


Activation Functions: Sigmoid vs. Softmax




Sigmoid,

$$\varphi(z) = \frac{1}{1 + e^{-z}}$$

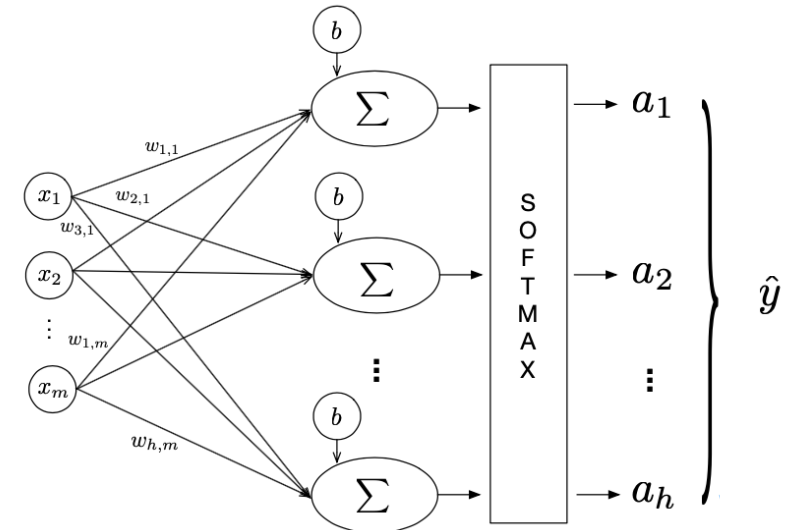
$$\hat{y} := \begin{cases} 1 & \text{if } \sigma(z) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$



?





?



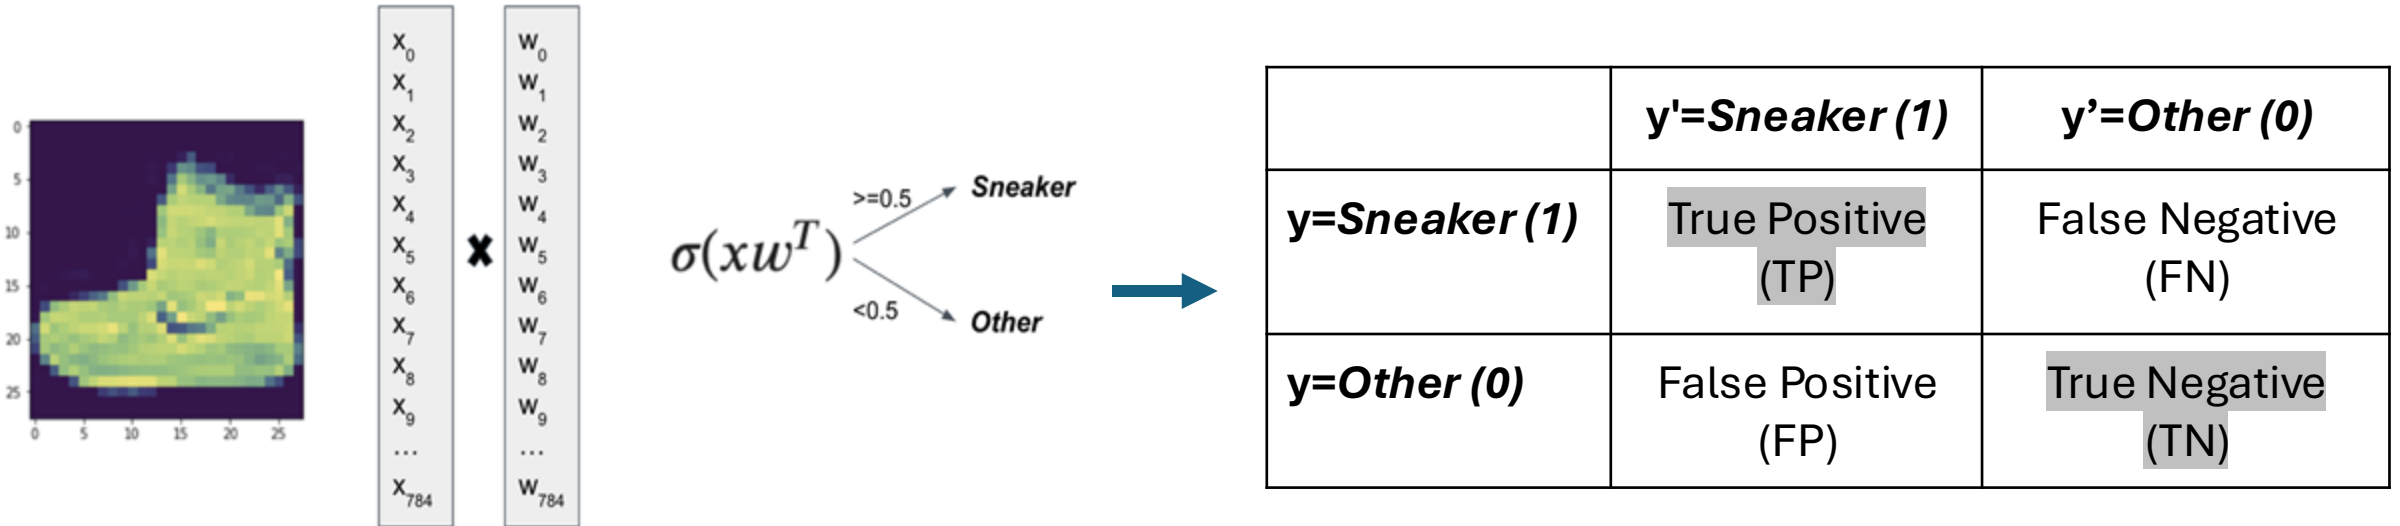
Logistic Regression: Accuracy

$$\text{Accuracy} = \frac{(TP+TN)}{(TP+TN+FP+FN)}$$

- A **single** threshold accuracy
 - Considers model quality only at one point
 - Threshold: typically, **0.5**
- Compare: Model confidence
- Compare: Prediction **0.49** vs. **0.51** for $y=1$?



The Confusion Matrix: Binary Classifier



Example, consider:

- Goal: Want to detect sneaker
- Metric: Accuracy

Actual \ Predicted	Other
Other	990
Sneaker	10

- Correct Predictions: 990
- Total Predictions: 1000
- Accuracy: 99.00%

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}} \times 100$$

$$\text{Accuracy} = \frac{990}{1000} \times 100 = 99.00\%$$

ROC AUC: Binary Classifier

- Example: 4 predictions
 - Class 1 = x2 (1s)
 - Class 0 = x2 (0s)

y	\hat{y}
0	0.3
1	0.5
0	0.6
1	0.9

Typically, 0.5 threshold

$$\hat{y} := \begin{cases} 1 & \text{if } \sigma(z) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

ROC AUC: Binary Classifier

- Example: 4 predictions
 - Class 1 = 2 (1s)
 - Class 0 = 2 (0s)
- Consider next all predictions \hat{y} as threshold

Typically, **0.5** threshold

y	\hat{y}	\hat{y}	0.3 threshold	$\leq 0.3: 0$ $> 0.3: 1$	0.5 threshold	0.6 threshold
0	0.3	0.3	0 (TN)		0 (TN)	0 (TN)
1	0.5	0.5	1 (TP)		0 (FN)	0 (FN)
0	0.6	0.6	1 (FP)		1 (FP)	0 (TN)
1	0.9	0.9	1 (TP)		1 (TP)	1 (TP)

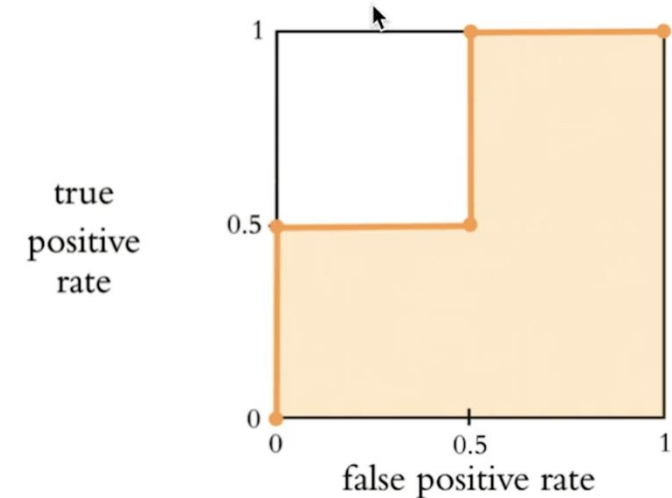
ROC AUC: Binary Classifier

- Example: 4 predictions
 - Class 1 = 2 (1s)
 - Class 0 = 2 (0s)
- Consider next all predictions y' as threshold

	y	\hat{y}		\hat{y}	0.3 threshold <small>$\leq 0.3 : 0$ $> 0.3 : 1$</small>		0.5 threshold		0.6 threshold
Typically, 0.5 threshold	0	0.3	→	0.3	0 (TN)	→	0 (TN)		0 (TN)
	1	0.5		0.5	1 (TP)		0 (FN)		0 (FN)
	0	0.6		0.6	1 (FP)		1 (FP)		0 (TN)
	1	0.9		0.9	1 (TP)		1 (TP)		1 (TP)
<div> <div> True Positive Rate = $\frac{TP}{TP+FN}$ False Positive Rate = $\frac{FP}{FP+TN}$ </div> <div> $\frac{2}{2+0} = 1.0$ $\frac{1}{1+1} = 0.5$ </div> <div> $\frac{1}{1+1} = 0.5$ $\frac{1}{1+1} = 0.5$ </div> <div> $\frac{1}{1+1} = 0.5$ $\frac{0}{0+2} = 0.0$ </div> </div>									

ROC AUC: Binary Classifier

- Example: 4 predictions
 - Class 1 = 2 (1s)
 - Class 0 = 2 (0s)
- Consider next all predictions y' as threshold



	y	\hat{y}	\hat{y}	0.3 threshold $\leq 0.3: 0$ $> 0.3: 1$	0.5 threshold	0.6 threshold
Typically, 0.5 threshold	0	0.3	0.3	0 (TN)	0 (TN)	0 (TN)
	1	0.5	0.5	1 (TP)	0 (FN)	0 (FN)
	0	0.6	0.6	1 (FP)	1 (FP)	0 (TN)
	1	0.9	0.9	1 (TP)	1 (TP)	1 (TP)

$$\text{True Positive Rate} = \frac{TP}{TP + FN}$$

$$\text{False Positive Rate} = \frac{FP}{FP + TN}$$

$$\frac{2}{2+0} = 1.0$$

$$\frac{1}{1+1} = 0.5$$

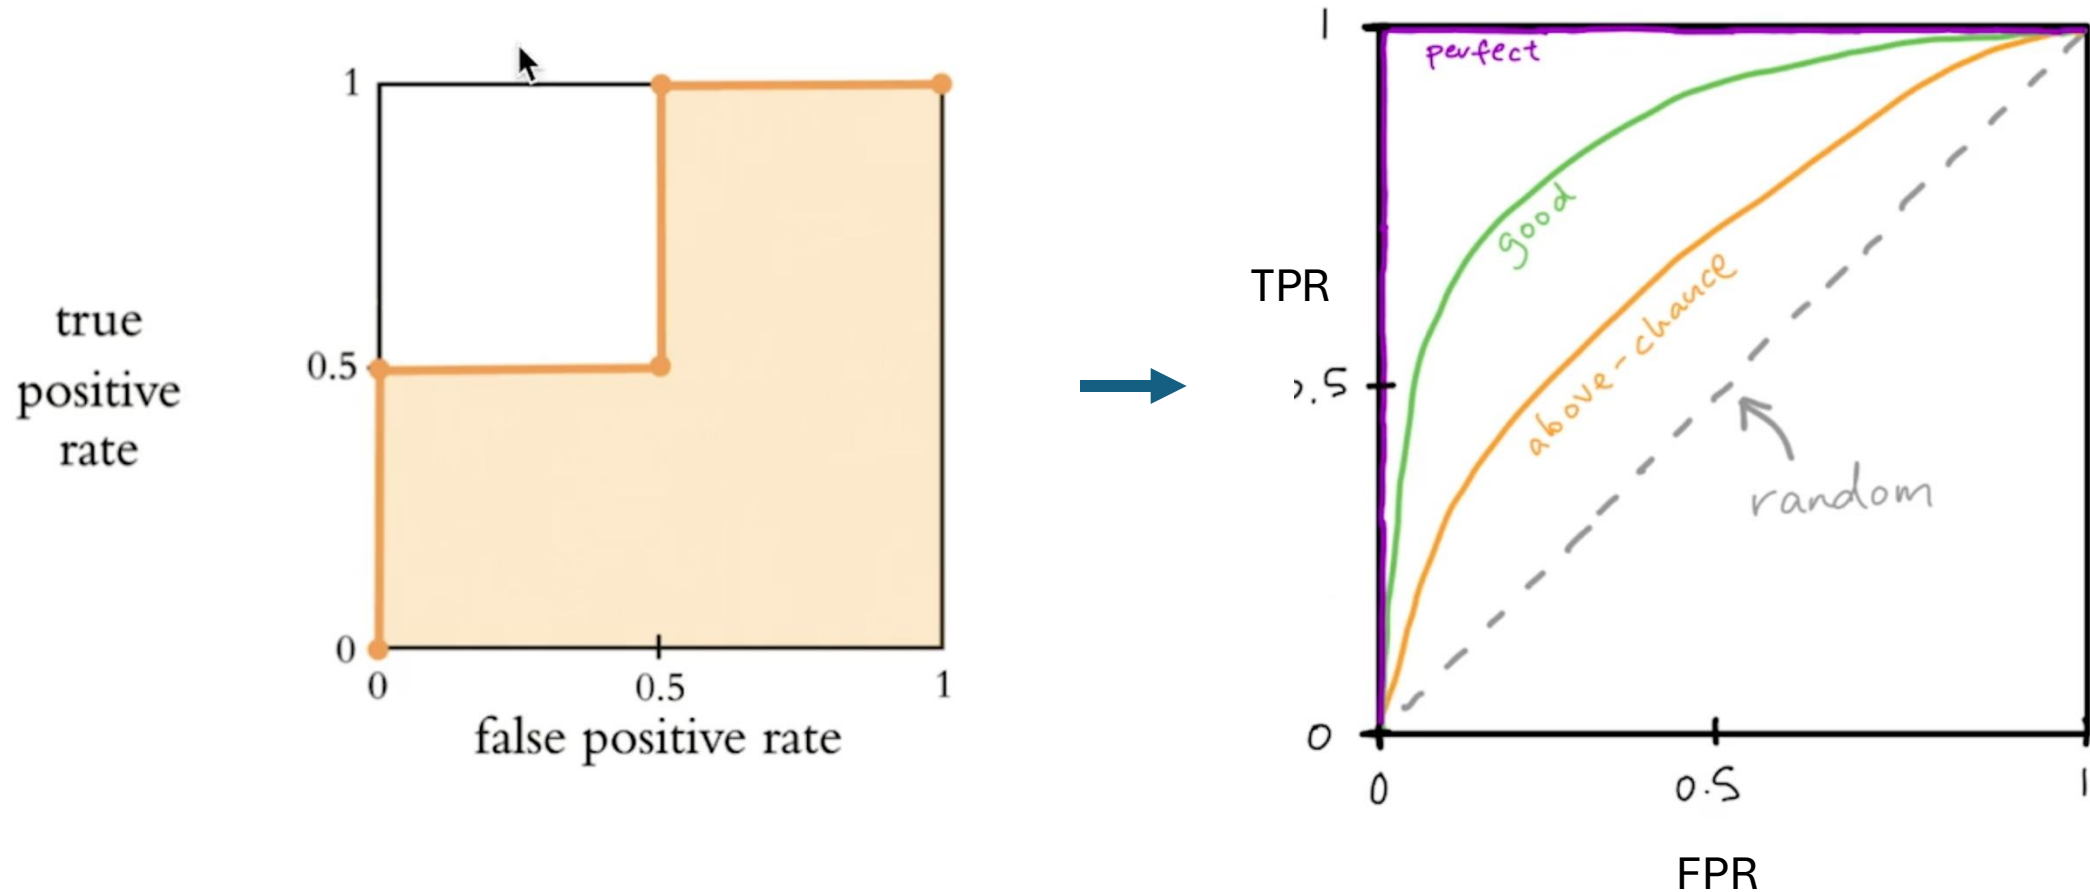
$$\frac{1}{1+1} = 0.5$$

$$\frac{1}{1+1} = 0.5$$

$$\frac{1}{1+1} = 0.5$$

$$\frac{0}{0+2} = 0.0$$

ROC AUC: Binary Classifier



Multiclass Confusion Matrix

- Classifier over multiple classes
 - Ex.: shirt-shirt = 514 correct predictions of shirt
 - Ex.: shirt-tshirt = 108 times predicted shirt as tshirt
- Multiclass ROC
 - [Scikit-learn: Multiclass ROC](#)
- Confusion Matrix:
 - [Sklearn: Confusion Matrix](#)

	$y' = 1$	$y' = 0$
$y = 1$	TP	FN
$y = 0$	FP	TN



True Label	t-shirt	787	6	12	59	6	1	108	0	20	1
	trouser	6	930	16	39	6	0	1	0	2	0
	pullover	15	0	669	8	188	1	101	0	18	0
	dress	32	18	7	849	42	1	47	0	4	0
	coat	0	3	93	39	771	0	88	0	6	0
	sandal	1	0	0	2	0	811	0	103	10	73
	shirt	148	3	133	36	131	1	514	0	34	0
	sneaker	0	0	0	0	0	36	0	898	0	66
	bag	2	1	12	11	3	6	24	9	931	1
	ankle boot	0	0	0	0	0	12	1	54	1	932
	t-shirt	trouser	pullover	dress	coat	sandal	shirt	sneaker	bag	ankle boot	
Predicted Label											

Precision, Recall, F1

emphasize correctness if we predict outcome is xyz (cost: a high num of FN)


- **Precision** is the percentage of **predicted** positives that were correctly classified

- how good a model is at predicting the positive class
- concerned with accuracy of the positive predictions
- increasing precision reduces recall and vice versa

optimizing for recall helps with minimizing the chance of not detecting xyz

- **Recall** is the percentage of **actual** positives that were correctly classified

- Consider when: Class imbalance



Consider Use-Cases:

- Fraud detection
- Subscribe or not to a product (Customer Acquisition)
- Medical Field: detection of a tumor with xyz imaging

Precision

$TP / (TP+FP)$

Recall

$TP / (TP+FN)$

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{2TP}{2TP + FN + FP}$$

Logistic Regression: Softmax Regression

$$z = w_1 x_1 + \dots + w_m x_m + b = \sum_{l=1}^m w_l x_l + b = \mathbf{w}^T \mathbf{x} + b$$

Sigmoid


Softmax

logit($p(y = 1 \mid \mathbf{x})$) = z

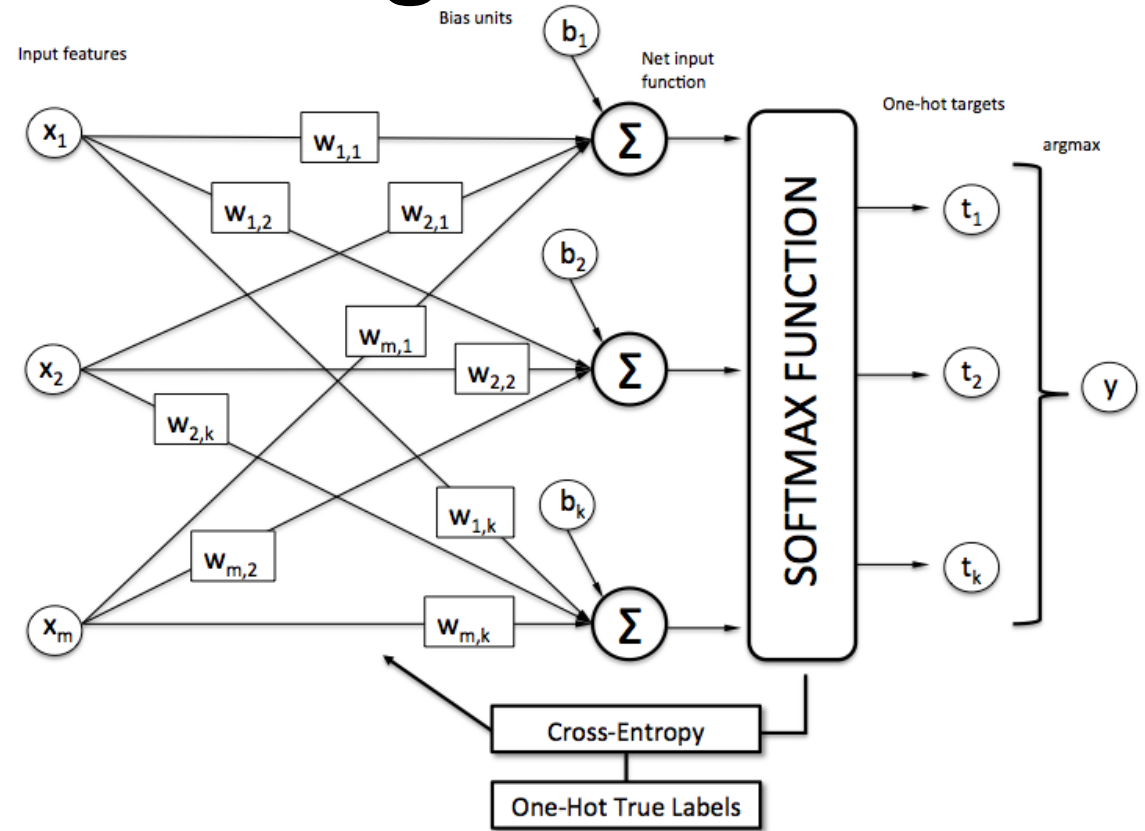
Inverse

$$\varphi(z) = \frac{1}{1 + e^{-z}}$$

$$P(y = j \mid \mathbf{z}^{(i)}) = \varphi_{\text{softmax}}(\mathbf{z}^{(i)}) = \frac{e^{z_j^{(i)}}}{\sum_{k=0}^K e^{z_k^{(i)}}}$$



- Another example: walkthrough by Rashka: [Multinomial Logistic](#)



$$\sigma_{\text{softmax}}(z_t^{[i]}) = \frac{e^{z_t^{[i]}}}{\sum_{j=1}^h e^{z_j^{[i]}}}$$

one-hot encode labels

$$\mathcal{L} = \sum_{i=1}^n \sum_{j=1}^h -y_j^{[i]} \log(a_j^{[i]})$$

over all training examples

h different class labels

Loss: Categorical Cross-Entropy

- Categorical cross-entropy
 - minimizes the distance between the probability distributions output by the model and the true distribution of the targets
- Handling labels in multiclass classification:
 - Encoding via categorical encoding (also known as one-hot encoding)
 - Use: `categorical_crossentropy` as a `loss` function
 - Encoding the labels as integers
 - Use: `sparse_categorical_crossentropy` `loss` function