

# DATASCI207-005/007

# Applied Machine Learning

Vilena Livinsky, PhD(c)

School of Information, UC Berkeley

Week 6: 02/12/2024 – 02/13/2024

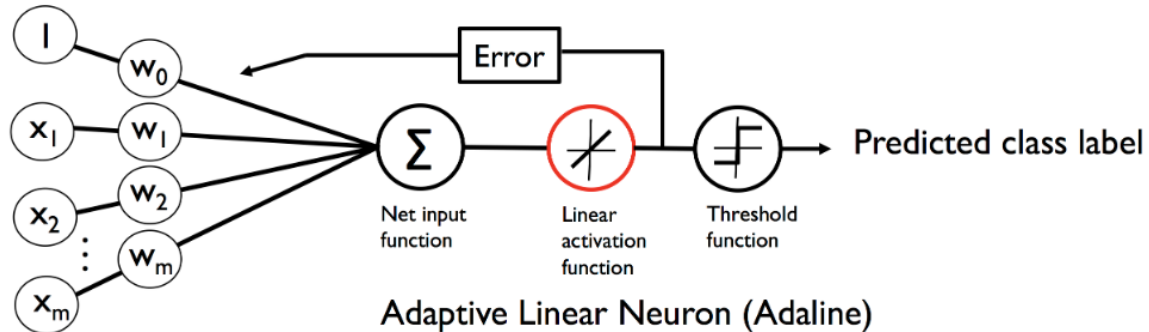
# Today's Agenda

- Feedforward Neural Nets
- Activation Functions
- Regularization
- Walkthroughs:
  - Regularization: L1, L2, Dropout (NNs)
  - XOR
    - TensorFlow: NNs




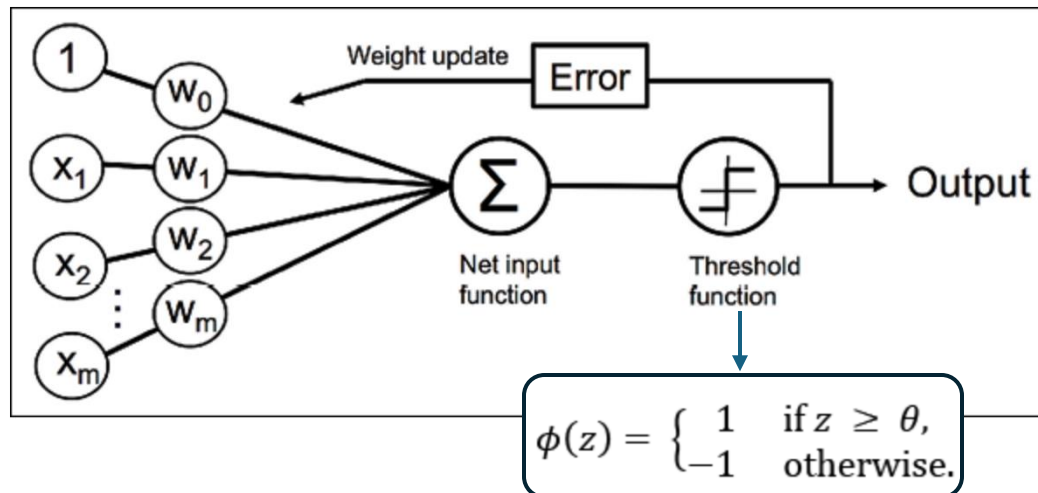
# Perceptron: Binary Classification

Review




linear activation function,  $\phi(z)$ ,


$$\phi(\mathbf{w}^T \mathbf{x}) = \mathbf{w}^T \mathbf{x}$$



# Constructing XOR



$x_1$	$x_2$	$h_1 = \text{OR}(x_1, x_2)$	$h_2 = \text{NAND}(x_1, x_2)$	$y = \text{AND}(h_1, h_2)$
0	0	0	1	0
0	1	1	1	1
1	0	1	1	1
1	1	1	0	0



# Neural Nets: Activation Functions & XOR

Animations of ReLU in action:

<https://blog.dailydoseofds.com/p/a-visual-and-intuitive-guide-to-what>

<https://nnfs.io/mvp/>

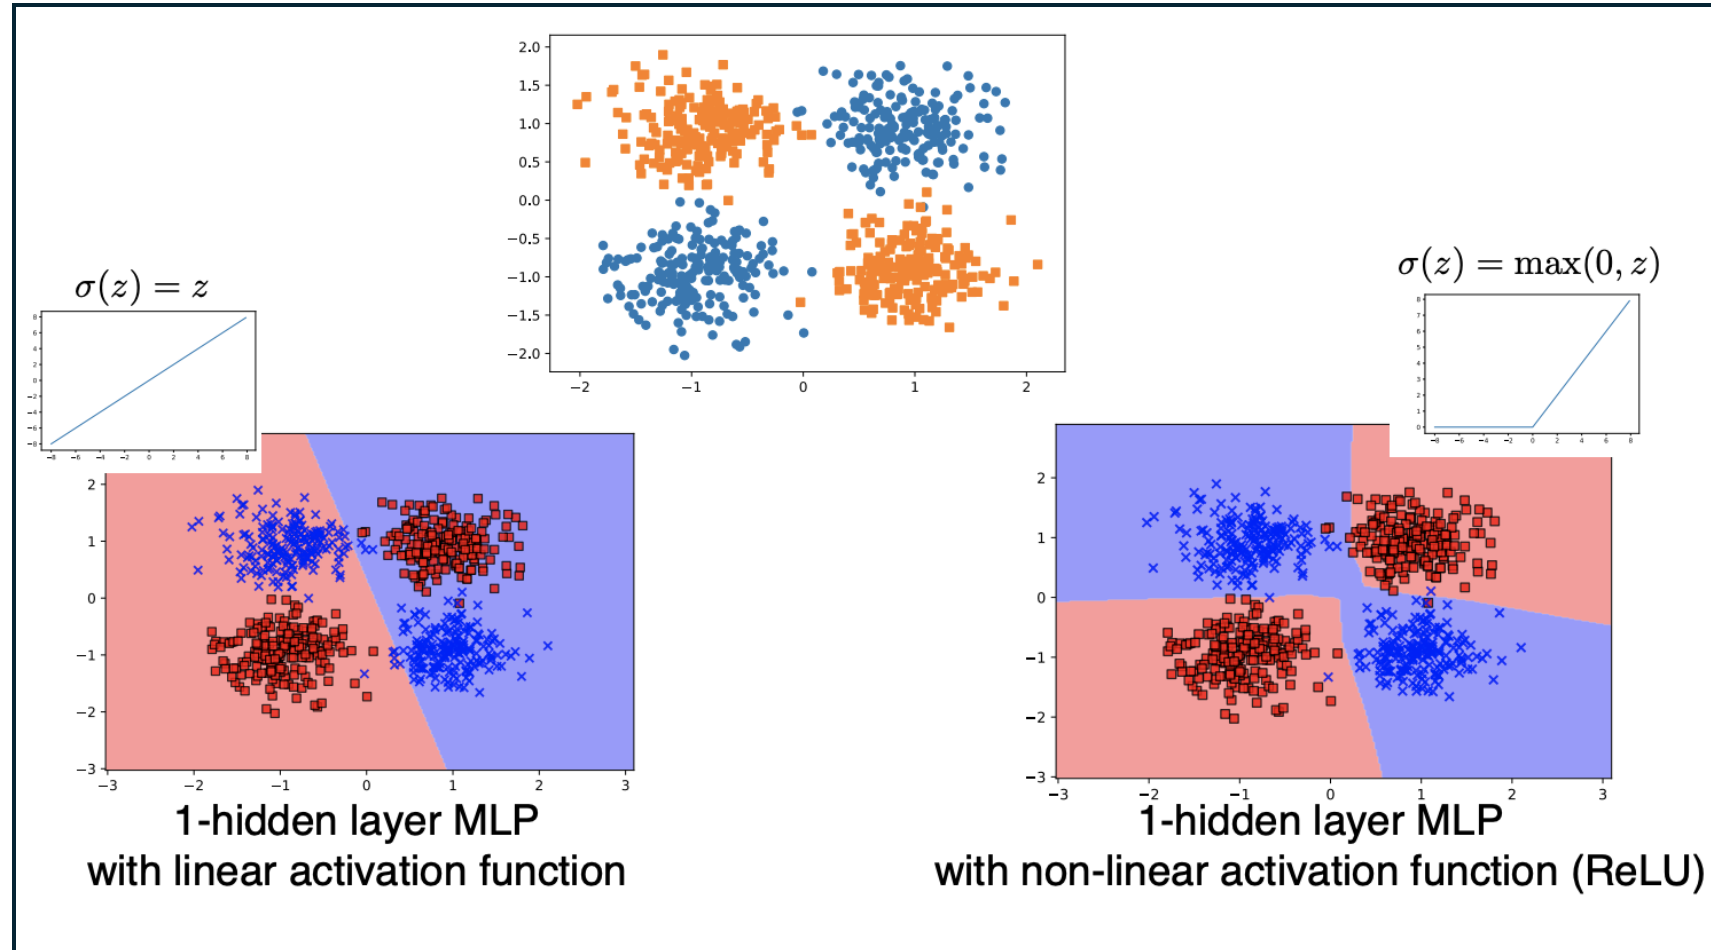


Image Source: Sebastian Raschka

<https://github.com/rasbt/stat453-deep-learning-ss21/blob/main/L09/code/xor-problem.ipynb>

# Multilayer Feedforward Neural Net (or MLP)

- Thus far:
- Ex.: Softmax Regression
  - Input – [Softmax – Output]
- Learning Algorithm
  - Gradient Descent

- Multilayer Perceptron (MLP)
  - Fully connected *feedforward* neural nets with 1(+) **hidden** layers
    - Feedforward: going in one direction, left to right (Input → Output)
    - More hyperparameters:
      - *number of layers*
      - *units*
- Learning Algorithm
  - **Backpropagation**
    - Gradient descent using the chain rule

# MLP: Examples, Classifiers

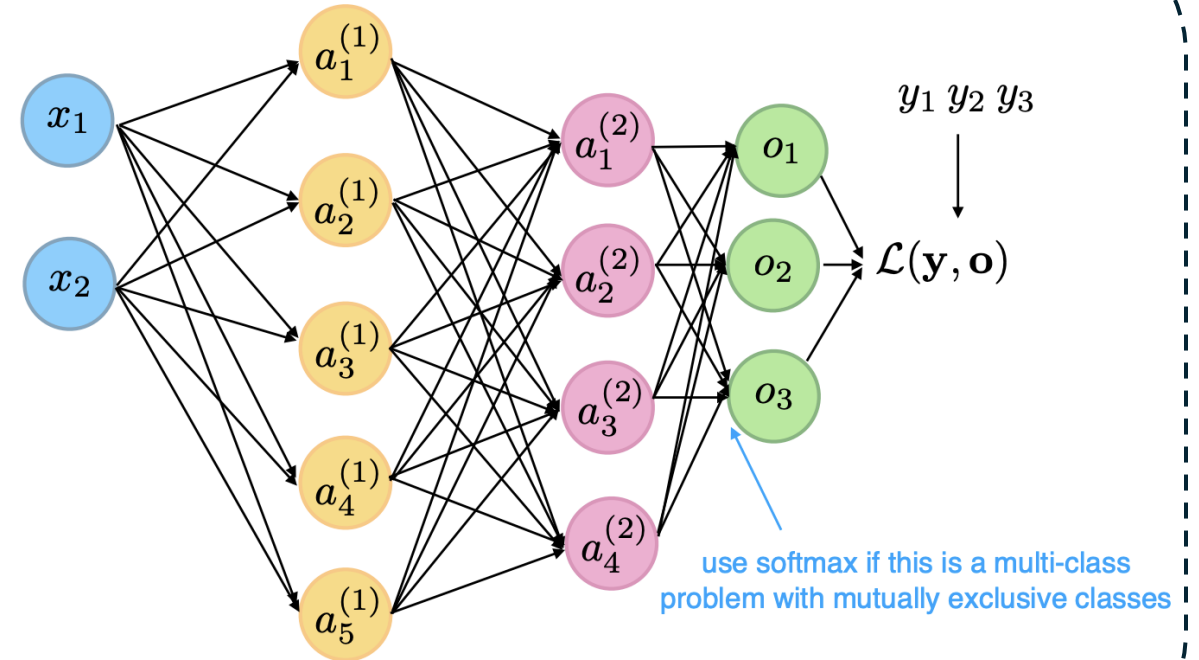
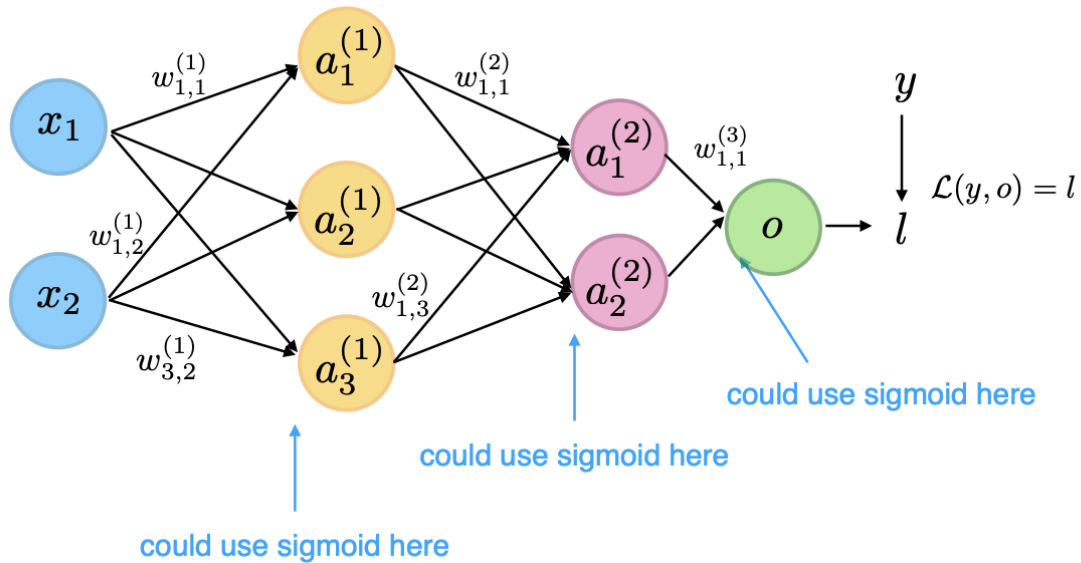


Image Source: Sebastian Raschka, Introduction to Deep Learning and Generative Models

# Activation Functions






Image Source:  
<http://deeplearning.stanford.edu/tutorial/supervised/MultiLayerNeuralNetworks/>

## Derivatives

### Sigmoid




### ReLU






# Activation Functions: Some Others


Identity




(Logistic) Sigmoid




ReLU (Rectified Linear Unit)



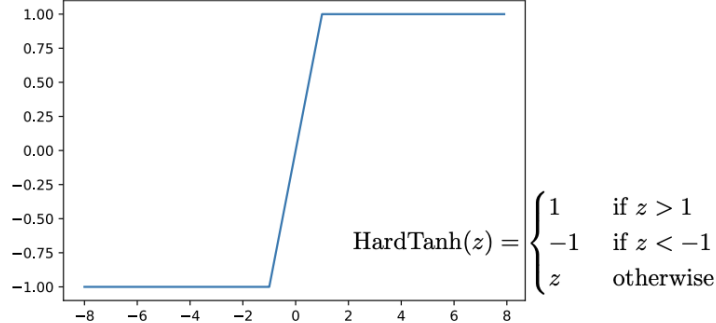
Leaky ReLU




Tanh ("tanH")



Hard Tanh



ELU (Exponential Linear Unit)



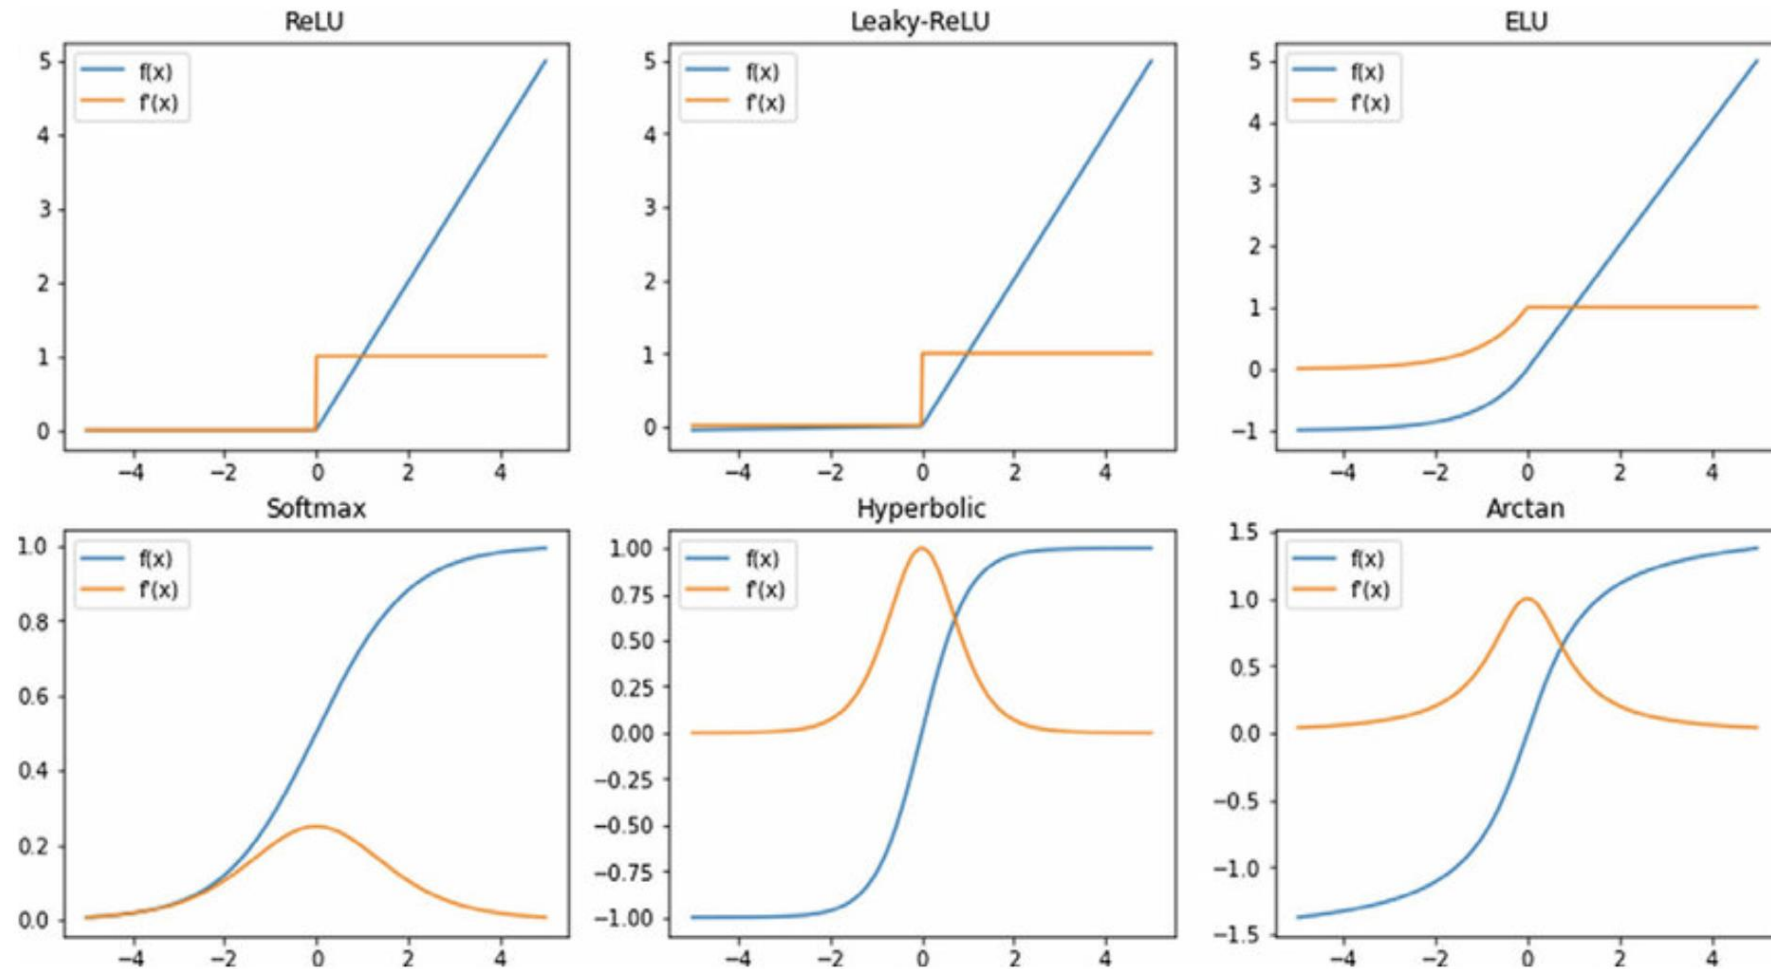
PReLU (Parameterized Rectified Linear Unit)

here, alpha is a trainable parameter

$$\text{PReLU}(z) = \begin{cases} z, & \text{if } z \geq 0 \\ \alpha z, & \text{otherwise} \end{cases}$$

$$\text{PReLU}(z) = \max(0, z) + \alpha \times \min(0, z)$$


# Activation Functions: Derivatives



Commonly used activation functions,  $f(x)$  and their derivatives,  $f'(x)$  for  $x \in [-5, 5]$

# Regularization: L1 (Lasso) vs. L2 (Ridge)


$$L1: \|\mathbf{w}\|_1 = \sum_{j=1}^m |w_j|$$



$$L2: \|\mathbf{w}\|_2^2 = \sum_{i=1}^m w_j^2$$

$$\text{Cost}_{\mathbf{w}, \mathbf{b}} = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{[i]}, \hat{y}^{[i]})$$

$$\text{L2-Regularized-Cost}_{\mathbf{w}, \mathbf{b}} = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{[i]}, \hat{y}^{[i]}) + \left[ \frac{\lambda}{n} \sum_j w_j^2 \right]$$




Keras API info on regularizers:

<https://keras.io/api/layers/regularizers/>

Example Implementation:

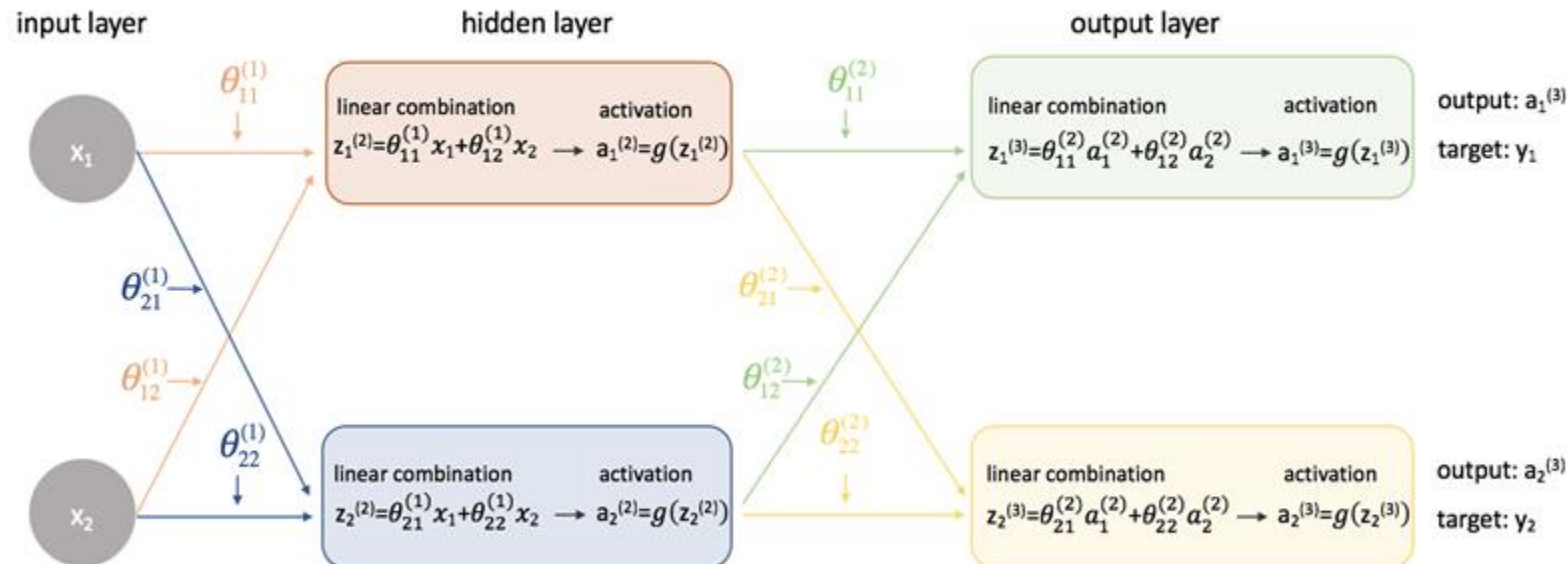
[https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first\\_edition/4.4-overfitting-and-underfitting.ipynb](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/first_edition/4.4-overfitting-and-underfitting.ipynb)

# Neural Nets: Backpropagation




$$\frac{\partial J(\theta)}{\partial \theta_2} = \left( \frac{\partial J(\theta)}{\partial a^{(3)}} \right) \left( \frac{\partial a^{(3)}}{\partial z} \right) \left( \frac{\partial z}{\partial \theta_2} \right)$$

$$\frac{\partial J(\theta)}{\partial \theta_1} = \left( \frac{\partial J(\theta)}{\partial a^{(3)}} \right) \left( \frac{\partial a^{(3)}}{\partial z^{(3)}} \right) \left( \frac{\partial z^{(3)}}{\partial a^{(2)}} \right) \left( \frac{\partial a^{(2)}}{\partial z^{(2)}} \right) \left( \frac{\partial z^{(2)}}{\partial \theta_1} \right)$$




## Non-linear Functions

If  $f$  is a linear function and  $g$  is a linear function, then their composition  $f(g(x))$  is a linear function



Without non-linear activations, all neural network functions could be reduced to a single layer

# Multilayer NNs: Cost Functions



Ref.: Raschka, S., & Mirjalili, V. (2019). Python Machine Learning, Third Edit.




Image Source: Li, H., Xu, Z., Taylor, G., Studer, C., & Goldstein, T. (2018). Visualizing the loss landscape of neural nets. Advances in neural information processing systems, 31.