

# DATASCI207-007

# Applied Machine Learning

Vilena Livinsky, PhD(c)

School of Information, UC Berkeley

Week 1: 01/09/2025

# Introductions: Welcome to the class!

---

- Current:
  - Lecturer in the School of Information at UC Berkeley
  - Full-time AI Data scientist for a leading U.S. financial services institution
    - Risk modelling
    - NLP/LLM modelling
    - Cybersecurity modelling
  - Researcher in speech science
    - PhD Candidate in Speech Science (CUNY)
- Fun fact:
  - Love travelling
  - Love very hot weather!



# About You



Undergraduate/ Graduate Major



Current position (if any)/ Would like to transition to?



DATASCI207: Topic most interested in exploring in this class?



Something you'd like to share about yourself?

A large orange circle occupies the left side of the slide, partially cut off by the edge.

# Live Sessions

- Each session is 90 minutes
  - Review of/ further deep dive into the week's material
  - Code demonstration/ Breakout room exercises
  - Q&A about the topic of the week



# Course Resources

## Async material

- in bCourses: <https://bcourses.berkeley.edu/>

## Live session

- Live session material will be posted on Slack
- Live lecture recording will be synced into each calendar invite in bCourses

## Slack Channels

- General Slack Channel: #datasci-207-2025-spring
- Section Slack Channel: #datasci-207-007-w25

## Schedule



- Live session: Th, 6:30 pm - 7:59PM PST
- Office hour: Wed, 8:00 - 9:00AM PST

# Today's Agenda

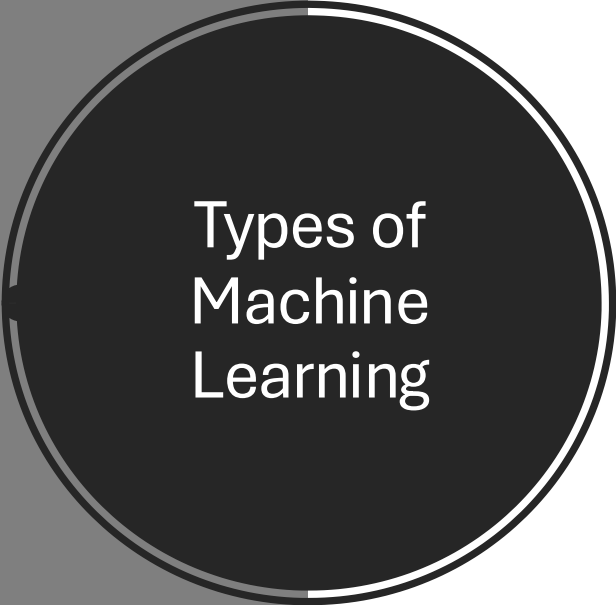
- Introductions
- General concepts of Machine Learning
  - Typical workflow for ML predictive modeling
- Walkthrough
  - NumPy Arrays
    - 1-D, 2-D, and n-D arrays (tensors)
- Walkthrough
  - a model build

# AI | Machine Learning | Deep Learning

- AI: automation of human intellectual tasks



- ML: ML learning system is trained (learning a function) given some data
  - Requires:
    - Input (exposure to examples)
    - Output
    - Metric for success (algo's current output vs expected output = feedback)
- Deep Learning: “layered representations learning”
  - Learning successive layers of increasingly meaningful representations (of data)
    - Neural network models



## Types of Machine Learning

- What types of Machine Learning come to mind?



# Types of Machine Learning

## Supervised Learning

- > Labeled data
- > Direct feedback
- > Predict outcome/future

## Unsupervised Learning

- > No labels
- > No feedback
- > Find hidden structure in data

## Reinforcement Learning

- > Decision process
- > Reward system
- > Learn series of actions

# Data & Learning a Function (Model)

## Supervised Learning




- > Labeled data
- > Direct feedback
- > Predict outcome/future

- Regression
  - Predicting continuous outcomes
- Classification
  - Predicting class labels

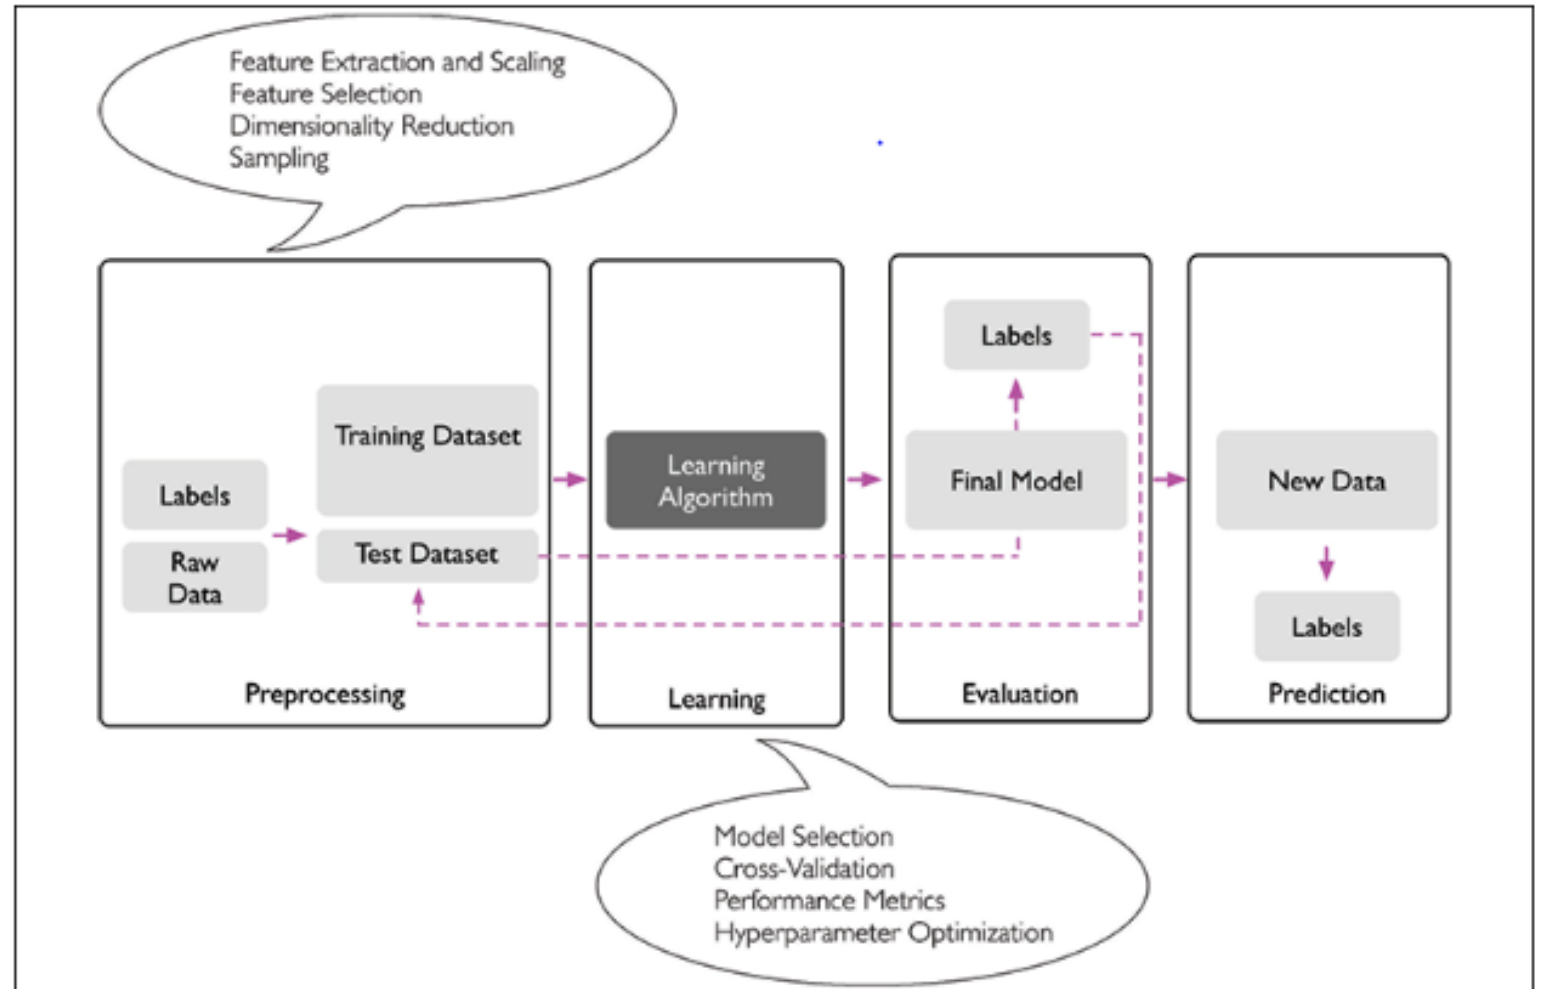
## Unsupervised Learning

- > No labels
- > No feedback
- > Find hidden structure in data

- Clustering
  - Finding subgroups in data



# Machine Learning Workflow



Q. Why do we need a train-test split?

# Tools

- ML/ AI APIs:
- CPU = general purpose processors
- GPUs/ TPUs = optimized to accelerate ML

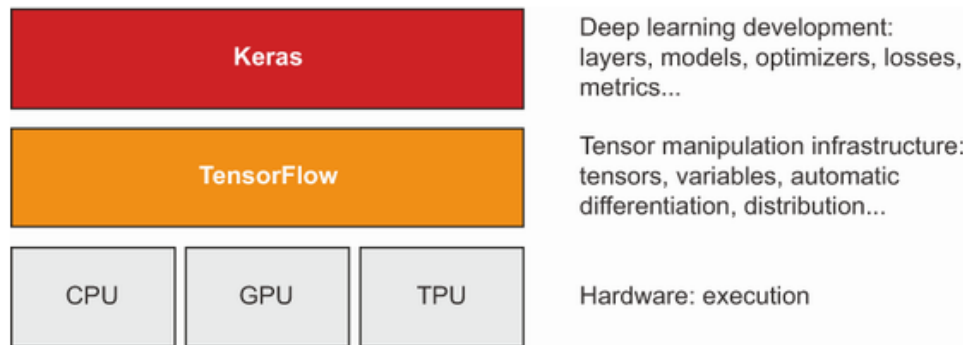


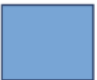
Figure 3.1 Keras and TensorFlow: TensorFlow is a low-level tensor computing platform, and Keras is a high-level deep learning API


- scikit-learn
  - will be used occasionally, especially for data preprocessing and model evaluation
- Other APIs:
  1. NumPy:
    - adds support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these arrays
  2. Pandas:
    - useful for data manipulation and analysis. In particular, it offers data structures and operations for manipulating numerical tables
  3. Matplotlib:
    - a plotting library

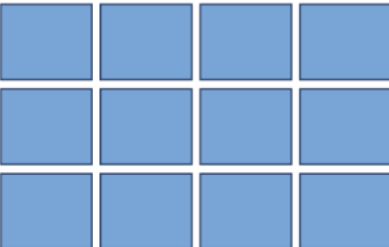
# Walkthrough

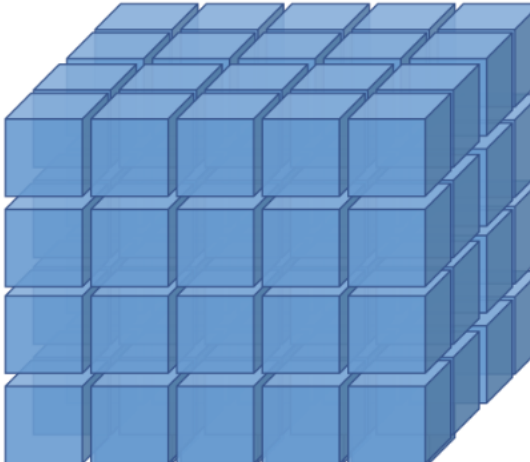
01a\_Introduction.ipynb

# Tensors

Rank 0:   
(scalar)

Rank 1:   
(vector)

Rank 2: (matrix)  


Rank 3: 

Tensors are *multi-dimensional* arrays with a *uniform type* (called a dtype).

- tensors are (kind of) like np.arrays

"Scalar" or "rank-0" tensor .

A scalar contains a single value, and no axes.

A "vector" or "rank-1" tensor is like a list of values. A vector has one axis.

A "matrix" or "rank-2" tensor has two axes.

TensorFlow

- <https://www.tensorflow.org/guide/tensor>

# Tensors: Note!

```
>>> x = np.array([12, 3, 6, 14, 7])
>>> x
array([12, 3, 6, 14, 7])
>>> x.ndim
1
```


This vector has five entries and so is called a *5-dimensional vector*. Don't confuse a 5D vector with a 5D tensor! A 5D vector has only one axis and has five dimensions along its axis, whereas a 5D tensor has five axes (and may have any number of dimensions along each axis). *Dimensionality* can denote either the number of entries along a specific axis (as in the case of our 5D vector) or the number of axes in a tensor (such as a 5D tensor), which can be confusing at times. In the latter case, it's technically more correct to talk about a *tensor of rank 5* (the rank of a tensor being the number of axes), but the ambiguous notation *5D tensor* is common regardless.

# Tensors


A scalar, shape: `[]`

4

A vector, shape: `[3]`






A matrix, shape: `[3, 2]`



There are many ways you might visualize a tensor with more than two axes.

A 3-axis tensor, shape: `[3, 2, 5]`





# Tensors: Examples

- A batch of 128 *grayscale* images of size  $256 \times 256$  could be stored in a tensor of shape (128, 256, 256, 1)
- A batch of 128 *color* images could be stored in a tensor of shape (128, 256, 256, 3)




Figure 2.4 A rank-4 image data tensor

# Walkthrough

01b\_Framing.ipynb